

ECS 158 Final Project:Parallelizing R's Phylobase::ShortestPath()

Raymond S. Chan, Alicia Luu, Bryan Ng
raschan@ucdavis.edu, ajuu@ucdavis.edu, bng@ucdavis.edu

Abstract

This report attempts to parallelize CRAN's phylobase package's shortestPath function in RSnow, OpenMP and CUDA. The function takes in a phylogeneic tree and produces the shortest path between two nodes. The RSnow implementation is built on top of the existing code and parallelizes the descendants function. The OpenMP implementation took a similar approach by parallelizing both the ancestors and descendants C functions. The CUDA implementation took the brute force approach to the shortestPath problem. Overall, our attempt to parallelize the shortest path function was partially successful. The original version has shown that it is quite fast and efficient in finding the shortest path.

1 Introduction and Motivation

Alongside the rise of "big data" in the recent years, bioinformatics has gained considerable momentum. But a consistent issue remains: What do we do with all the data and how do we make sense of it at a reasonable rate? The R community has attempted to resolute these issues. For this report we are examining the R's phylobase package which provides the base class and functions for phylogenetic or evolutionary structures and comparative data. We will be focusing our efforts in making the "treewalk" utility functions, such as finding descendants/ancestors and shortest paths quickly through three different parallel programming models (RSnow, OpenMP and CUDA).

2 Approaches

Only pseudocode or key chunks of code of each implementation are shown or described, see Appendix A for code details.

2.1 Original

The original R implementation calculated the shortest path between the two nodes of interest by first calculating their Most Recent Common Ancestor (MRCA). The MRCA's descendants are then calculated and compared to the the two nodes's ancestors. Any overlap is stored and that is the shortest path.

The C version of the descendants function, called Cdescendant(), works by first marking a given node in a preordered list of edges. The direct descendants (or children) of the given node are then marked. Cdescendant() then iterates through the other edges and marks each marked node's direct descendants (children).

The C version of the ancestors function works the same as Cdescendants, but direct *ancestors* (or parents) are marked instead of direct *descendants* (children).

Pseudocode (source code in original phylobase package):

```

1 descendants(tree, given_node){
2   #let x be the edges of tree listed in PREORDER,
3   #with the older node occupying the first column
4   # C function call
5   isDescendant <- Cdescendants(x[,1], x[,2], given_node)
6   retval <- getNode(tree, isDescendant)
7 }
8
9 ancestors(tree, node1){
10  #let x be the edges of tree listed in POSTORDER,
11  #with the older node occupying the first column
12  # C function call
13  isAncestor <- Cancestors(x[,1], x[,2], given_node)
14  retval <- getNode(tree, isAncestor)
15 }
16
17 MRCA(tree, node1, node2 ... noden){
18   nodes <- unique(node1, node2, ..., noden)
19   ancests <- lapply(nodes, ancestors, phy=phy, type="ALL")
20   retval <- getNode(phy, max(Reduce(intersect, ancests)))
21 }
22 shortestPath(tree, node1, node2){
23   t1 <- getNode(tree, node1)
24   t2 <- getNode(tree, node2)
25
26   # most recent common ancestor
27   comAnc <- MRCA(tree, t1, t2)
28   desComAnc <- descendants(tree, comAnc)
29
30   # path: common ancestor to t1
31   ancT1 <- ancestors(x, t1)
32   path1 <- intersect(desComAnc, ancT1)
33
34   # path: common ancestor to t2
35   ancT2 <- ancestors(x, t2)
36   path2 <- intersect(desComAnc, ancT2)
37
38   # union of the path above paths
39   retval <- union(path1, path2)
40 }

```

2.2 RSnow

The RSnow implementation builds on top of original R version by parallelizing the descendants function. In order to make independent subproblems, for a given node, every other node keeps marching upwards to its ancestors until they either reached the root or encountered the given node. We were unable to parallelize the ancestor function. The original R version seems to have taken the most efficient serial approach for calculating a given node's ancestors.

Pseudocode: (see code in Appendix A.1)

```
1 descendants(tree, given_node){  
2   #let x be the list of all nodes except the given_node  
3   for i from 1 to height of tree  
4     if x == given_node -> append node to retval  
5     update x with its corresponding ancestor  
6   return retval  
7 }
```

2.3 OpenMP

OpenMP allows parallelism to be added without changing a significant amount of existing code. And that's exactly what we did implement to implement an OpenMP version of the original shortestPath program. We decided to parallelize the two serial C functions called and implied by the shortestPath function in treewalk.R. Those two C functions are ancestors() in ancestors.c and descendants() in descendants.c. Because serial for loops comprised the bulk of the ancestor and descendant C functions, the OpenMP implementation is embarrassingly parallel. The functions are also small so there isn't a need for explicit barriers.

See code in Appendix A.2.

2.4 CUDA

Our CUDA implementation of the shortestPath function utilizes the GPU to find all ancestors of a given pair of nodes and then construct the shortest path between them. Our implementation assumes CSIF's pc43's resources, which are 1024 threads per block and 1 GB of global memory. We assume the given data can fit in our GPUs global memory. This assumption may limit the test we will be able to perform. Our solution utilizes the fact that the shortest path between two nodes in a tree must converge at the lowest common ancestor of both nodes. In cases, where one node is an ancestor of another, the shortest path is then found by traversing the parents of the child node. We parallelized our code by finding both sets of ancestors of the given nodes at the same time. Since neither node needs to know about the other to find its own ancestors, this problem can be done independently of each other. Both sets of ancestors are then traversed to find the shortest path. We were unable to parallelize this part of the solution since each list of ancestors must be checked to find overlapping elements.

See code in Appendix A.3.

3 Experiment Results

These are the unprocessed results from using R's `system.time()` on the RSnow and OMP `shortestPath` functions on `testXXX.tre` test cases, where XXX denotes the numbers of leaves the tree has. The serial label denotes the original R implementation. We did not manage to get the CUDA implementation to correctly interface with R.

```
#msdir/set1/test10.tre
#serial
c(0.015, 0.011, 0.012, 0.012, 0.012, 0.011, 0.012, 0.008, 0.011, 0.012)
#RSnow
c(0.068, 0.017, 0.018, 0.017, 0.017, 0.017, 0.017, 0.018, 0.017, 0.019)
#OMP
c(0.016, 0.011, 0.012, 0.011, 0.011, 0.012, 0.011, 0.009, 0.012, 0.011)

#msdir/set1/test100.tre
#serial
c(0.027, 0.030, 0.022, 0.020, 0.020, 0.021, 0.021, 0.020, 0.021, 0.021)
#RSnow
c(0.155, 0.132, 0.117, 0.118, 0.130, 0.125, 0.136, 0.111, 0.115, 0.114)
#OMP
c(0.046, 0.025, 0.021, 0.021, 0.021, 0.020, 0.021, 0.021, 0.020, 0.021)

#msdir/set1/test1000.tre
#serial
#c(0.131, 0.128, 0.125, 0.122, 0.124, 0.136, 0.129, 0.123, 0.123, 0.126)
#RSnow
c(16.454, 17.481, 16.778, 17.686, 12.808, 16.032, 15.914, 16.816, 13.582, 14.389)
#OMP
#c(0.132, 0.130, 0.151, 0.124, 0.133, 0.134, 0.125, 0.124, 0.123, 0.126)

#msdir/set1/test10000.tre
#serial
c(1.819, 1.851, 1.850, 1.789, 1.754, 1.840, 1.809, 1.699, 1.735, 1.787)
#RSnow
c(0,0,0,0,0,0,0,0,0,0) #program crashed on this case
#OMP
c(1.783, 1.841, 1.820, 1.751, 1.805, 1.830, 1.796, 1.705, 1.733, 1.784)
```

4 Discussion

4.1 RSnow

The RSnow version was tested with four localhost worker nodes. The large majority of the runs were slower than the serial version. Looking back at our approach to the problem, it was quite demanding for resources, specifically memory hungry. It was not a good idea for have four localhost nodes tackle the issue.

4.2 OpenMP

The OpenMP version works marginally faster than the original code on large phylogenetic trees (approximately 10000 leaves). With small trees, serial "for" loops seem to be efficient in finding ancestor and descendant nodes.

4.3 CUDA

Compared to the serial version, the cuda implementation performed slower in most test cases. While the CUDA version can compute both given nodes ancestors at the same time, it must also load the entire tree into the GPUs memory.

5 Conclusion

In the attempt to parallelize the shortest path function in the phylobase package, the RSnow implementation illustrated the efficiency of the original code and its serial search algorithm. The OpenMP implementation showed the utility of multithreading for large data. The CUDA implementation could not be timed, because of R/C interafacing issues.

6 Acknowledgements

We would like to thank Professor Norman Matloff for his guidance and knowledge presented during lectures. This work is the result of a final project for ECS 158 Winter Quarter 2015. His open-source textbook, blog and various tutorial were an essential part of our learning. We would also like to thank the teaching assistance Shengren Li for offering invaluable advice and feedback on our codes (especically our CUDA) throughout the quarter.

7 Appendix

A Codes

The RSnow descendant functions calculates decendant in reverse by looking at ancestors.

A.1 RSnow Code

```
1 SNOW <- function(x, size, root, type=c("descendants")){
2   ans <- rep(0, size)
3   mystart <- (myid-1)*length(x)+1
4   myend <- myid*length(x)
5
6   type <- match.arg(type)
7   if (type == "descendants"){
8     v1 <- descendant
9     v2 <- ancestor
10    #initialization
11    temp <- v1[mystart:myend]
12
13    #second and beyond iteration
14    for (j in 1:size){
15      if (node %in% temp){
16        setthese <- which(temp == node) + mystart-1
17        ans[setthese] <- 1
18      }
19      blah <- rep(-1, length(temp))
20      for (i in (1:length(temp))){
21        matched_pos <- which(v1 == temp[i])
22        if (length(matched_pos) != 0){
23          blah[which(temp == temp[i])] <- matched_pos
24        }
25        else{#matched_pos == 0
26          ## R is 1 INDEXED!
27          if (type == "descendants"){
28            blah[i] <- 1
29          }
30        }
31      }#for i
32      #go to your parents set
33      difference <- length(temp) - length(v2[blah])
34      temp <- v2[blah]
35      if (difference > 0){
36        temp <- c(rep(0, difference), temp)
37      }
38      if (node %in% temp){
39        setthese <- which(temp == node) + mystart-1
40        ans[setthese] <- 1
41      }
42    }#j loop
43  }#new endif for type==descendants
44  return(ans)
45 }# end SNOW
46
47 setmyid <- function(i){
48   myid <- i
49 }
```

```

50
51 ## get descendants with RSnow
52 RSnowdescendants <- function (phy, node, type=c("tips", "children", "all"), cls) {
53   type <- match.arg(type)
54
55   ## look up nodes, warning about and excluding invalid nodes
56   oNode <- node
57   node <- getNode(phy, node, missing="warn")
58   isValid <- !is.na(node)
59   node <- as.integer(node[isValid])
60
61   if (type == "children") {
62     res <- lapply(node, function(x) children(phy, x))
63     ## if just a single node, return as a single vector
64     if (length(res)==1) res <- res[[1]]
65   } else {
66     ## edge matrix must be in preorder for the C function!
67     #if (phy@order=="preorder") {
68       edge <- phy@edge
69     #} else {
70     #   edge <- reorder(phy, order="postorder")@edge
71     #}
72     ## extract edge columns
73     ancestor <- as.integer(edge[, 1])
74     descendant <- as.integer(edge[, 2])
75
76     ## return indicator matrix of ALL descendants (including self)
77     #isDes <- .Call("descendants", node, ancestor, descendant)
78     clusterExport(cls, c("node", "ancestor", "descendant", "setmyid", "SNOW"),
79                   envir=environment())
79     dexgrps <- splitIndices(length(ancestor), length(cls))
80     rootdex <- which(phy@edge[,1] == 0)
81     clusterApply(cls, 1:length(cls), setmyid)
82     newisDes <- clusterApply(cls, dexgrps, SNOW, length(ancestor), rootdex,
83                             "
84                             descendants")
85     isDes <- (matrix(Reduce('+', newisDes), nrow=length(ancestor), ncol=1))
86     storage.mode(isDes) <- "logical"
87     ## for internal nodes only, drop self (not sure why this rule?)
88     int.node <- intersect(node, nodeId(phy, "internal"))
89     isDes[cbind(match(int.node, descendant),
90                match(int.node, node))] <- FALSE
91
92     ## if only tips desired, drop internal nodes
93     if (type=="tips") {
94       isDes[descendant %in% nodeId(phy, "internal"),] <- FALSE
95     }
96     ## res <- lapply(seq_along(node), function(n) getNode(phy,
97     ##   descendant[isDes[,n]]))
98     res <- getNode(phy, descendant[isDes[, seq_along(node)]])
99   }
100   ## names(res) <- as.character(oNode[isValid])
101   res
102 }
103 #####
104 # shortestPath
105 #####
106

```

```

107 RSnowshortestPath <- function(phy, node1, node2, cls){
108
109     ## conversion from phylo, phylo4 and phylo4d
110     if (class(phy) == "phylo4d") {
111         x <- extractTree(phy)
112     }
113     else if (class(phy) != "phylo4"){
114         x <- as(phy, "phylo4")
115     }
116     ## some checks
117     t1 <- getNode(x, node1)
118     t2 <- getNode(x, node2)
119     if(any(is.na(c(t1,t2)))) stop("wrong node specified")
120     if(t1==t2) return(NULL)
121
122     ## main computations
123     comAnc <- MRCA(x, t1, t2) # common ancestor
124     desComAnc <- RSnowdescendants(x, comAnc, type="all", cls)
125     ancT1 <- ancestors(x, t1, type="all")
126     path1 <- intersect(desComAnc, ancT1) # path: common anc -> t1
127
128     ancT2 <- ancestors(x, t2, type="all")
129     path2 <- intersect(desComAnc, ancT2) # path: common anc -> t2
130
131     res <- union(path1, path2) # union of the path
132     ## add the common ancestor if it differs from t1 or t2
133     if(!comAnc %in% c(t1,t2)){
134         res <- c(comAnc, res)
135     }
136
137     res <- getNode(x, res)
138
139     return(res)
140 } # end shortestPath

```

A.2 OpenMP Code

As stated approve in the approach section, very few lines were changed. OMP for's were added everywhere. The R code was modified so it would not interfere with the original code syntactically.

```

1
2 #include <R.h>
3 #include <Rinternals.h>
4 #include "omp.h"
5 int num_threads=8;
6 void omp_set_num_threads(int num_threads);
7
8 SEXP OMPancestors(SEXP nod, SEXP anc, SEXP des) {
9
10     int numEdges = length(anc);
11     int numNodes = length(nod);
12
13     int* nodes = INTEGER(nod);
14     int* ancestor = INTEGER(anc);
15     int* descendant = INTEGER(des);
16
17     int parent = 0;
18     SEXP isAncestor;
19

```



```

20     PROTECT(isAncestor = allocMatrix(INTSXP, numEdges, numNodes));
21 #pragma omp parallel for collapse(2)
22     for (int n=0; n<numNodes; n++) {
23         //pragma omp parallel for
24         for (int i=0; i<numEdges; i++) {
25             if (nodes[n]==descendant[i]) {
26                 INTEGER(isAncestor)[i + n*numEdges] = 1;
27             } else {
28                 INTEGER(isAncestor)[i + n*numEdges] = 0;
29             }
30         }
31     }
32
33 #pragma omp parallel for collapse(2)
34     for (int n=0; n<numNodes; n++) {
35         //pragma omp parallel for
36         for (int i=0; i<numEdges; i++) {
37             if (INTEGER(isAncestor)[i + n*numEdges]==1) {
38                 parent = ancestor[i];
39                 for (int j=i+1; j<numEdges; j++) {
40                     if (descendant[j]==parent) {
41                         INTEGER(isAncestor)[j + n*numEdges]=1;
42                     }
43                 }
44             }
45         }
46     }
47
48     UNPROTECT(1);
49     return isAncestor;
50 }
51
52 SEXP OMPdescendants(SEXP nod, SEXP anc, SEXP des) {
53
54     int numEdges = length(anc);
55     int numNodes = length(nod);
56
57     int* nodes = INTEGER(nod);
58     int* ancestor = INTEGER(anc);
59     int* descendant = INTEGER(des);
60
61     int child = 0;
62     SEXP isDescendant;
63
64     PROTECT(isDescendant = allocMatrix(INTSXP, numEdges, numNodes));
65 #pragma omp parallel for collapse(2)
66     for (int n=0; n<numNodes; n++) {
67         for (int i=0; i<numEdges; i++) {
68             if (nodes[n]==descendant[i]) {
69                 INTEGER(isDescendant)[i + n*numEdges] = 1;
70             } else {
71                 INTEGER(isDescendant)[i + n*numEdges] = 0;
72             }
73         }
74     }
75
76 #pragma omp parallel for collapse(2)
77     for (int n=0; n<numNodes; n++) {
78         for (int i=0; i<numEdges; i++) {
79             if (INTEGER(isDescendant)[i + n*numEdges]==1) {
80                 child = descendant[i];
81                 for (int j=i+1; j<numEdges; j++) {
82                     if (ancestor[j]==child) {
83                         INTEGER(isDescendant)[j + n*numEdges] = 1;
84                     }
85                 }
86             }
87         }
88     }

```

```

88 }
89 UNPROTECT(1);
90 return isDescendant;
91 }

```

```

1  ## get descendants [recursively]
2  OMPdescendants <- function (phy, node, type=c("tips","children","all")) {
3      type <- match.arg(type)
4
5      ## look up nodes, warning about and excluding invalid nodes
6      oNode <- node
7      node <- getNode(phy, node, missing="warn")
8      isValid <- !is.na(node)
9      node <- as.integer(node[isValid])
10
11     if (type == "children") {
12         res <- lapply(node, function(x) children(phy, x))
13         ## if just a single node, return as a single vector
14         if (length(res)==1) res <- res[[1]]
15     } else {
16         ## edge matrix must be in preorder for the C function!
17         if (phy@order=="preorder") {
18             edge <- phy@edge
19         } else {
20             edge <- reorder(phy, order="preorder")@edge
21         }
22         ## extract edge columns
23         ancestor <- as.integer(edge[, 1])
24         descendant <- as.integer(edge[, 2])
25
26         ## TODO: REPLACE C call with OMP implementation of descendants
27         ## return indicator matrix of ALL descendants (including self)
28         isDes <- .Call("OMPdescendants", node, ancestor, descendant)
29         storage.mode(isDes) <- "integer"
30     }
31     OMPdescendants <- function (phy, node, type=c("tips","children","all")) {
32         type <- match.arg(type)
33
34         ## look up nodes, warning about and excluding invalid nodes
35         oNode <- node
36         node <- getNode(phy, node, missing="warn")
37         isValid <- !is.na(node)
38         node <- as.integer(node[isValid])
39
40         if (type == "children") {
41             res <- lapply(node, function(x) children(phy, x))
42             ## if just a single node, return as a single vector
43             if (length(res)==1) res <- res[[1]]
44         } else {
45             ## edge matrix must be in preorder for the C function!
46             if (phy@order=="preorder") {
47                 edge <- phy@edge
48             } else {
49                 edge <- reorder(phy, order="preorder")@edge
50             }
51             ## extract edge columns
52             ancestor <- as.integer(edge[, 1])
53             descendant <- as.integer(edge[, 2])
54
55             ## TODO: REPLACE C call with OMP implementation of descendants

```

```

55     ## return indicator matrix of ALL descendants (including self)
56     isDes <- .Call("OMPdescendants", node, ancestor, descendant)
57     storage.mode(isDes) <- "logical"
58
59     ## for internal nodes only, drop self (not sure why this rule?)
60     int.node <- intersect(node, nodeId(phy, "internal"))
61     isDes[cbind(match(int.node, descendant),
62                 match(int.node, node))] <- FALSE
63
64     ## if only tips desired, drop internal nodes
65     if (type=="tips") {
66         isDes[descendant %in% nodeId(phy, "internal"),] <- FALSE
67     }
68
69     res <- getNode(phy, descendant[isDes[, seq_along(node)]]))
70 }
71
72     res
73 }#OMPdescendants
74
75 ## get ancestors (all nodes)
76 OMPancestors <- function (phy, node, type=c("all","parent","ALL")) {
77     type <- match.arg(type)
78
79     ## look up nodes, warning about and excluding invalid nodes
80     oNode <- node
81     node <- getNode(phy, node, missing="warn")
82     isValid <- !is.na(node)
83     node <- as.integer(node[isValid])
84
85     if (length(node) == 0) {
86         return(NA)
87     }
88
89     if (type == "parent") {
90         res <- lapply(node, function(x) ancestor(phy, x))
91     } else {
92         ## edge matrix must be in postorder for the C function!
93         if (phy@order=="postorder") {
94             edge <- phy@edge
95         } else {
96             edge <- reorder(phy, order="postorder")@edge
97         }
98         ## extract edge columns
99         ancestor <- as.integer(edge[, 1])
100        descendant <- as.integer(edge[, 2])
101
102        ## TODO: REPLACE C call with OMP implementation of ancestors
103        ## return indicator matrix of ALL ancestors (including self)
104        isAnc <- .Call("OMPancestors", node, ancestor, descendant)
105        storage.mode(isAnc) <- "logical"
106
107        ## drop self if needed
108        if (type=="all") {
109            isAnc[cbind(match(node, descendant), seq_along(node))] <- FALSE
110        }
111        res <- lapply(seq_along(node), function(n) getNode(phy,
112            descendant[isAnc[,n]]))
113    }

```

```

114     names(res) <- as.character(oNode[isValid])
115
116     ## if just a single node, return as a single vector
117     if (length(res)==1) res <- res [[1]]
118     res
119 }#OMPancestors
120
121 OMPMRCA <- function(phy, ...) {
122     nodes <- list(...)
123     ## if length==1 and first element is a vector,
124     ## use it as the list
125     if (length(nodes)==1 && length(nodes[[1]])>1) {
126         nodes <- as.list(nodes[[1]])
127     }
128
129     ## Correct behavior when the root is part of the nodes
130     testNodes <- lapply(nodes, getNode, x=phy)
131     ## BMB: why lapply, not sapply?
132     lNodes <- unlist(testNodes)
133     if (any(is.na(lNodes)))
134         stop("nodes not found in tree: ", paste(names(lNodes)[is.na(lNodes)],
135                                                  collapse=", "))
136
137     uniqueNodes <- unique(testNodes)
138     root <- nTips(phy)+1
139     ## Handles case where root is a node of interest, return root
140     if(root %in% uniqueNodes) {
141         res <- getNode(phy, root)
142         return(res)
143     }
144     ## Correct behavior in case of MRCA of identical taxa
145     if(length(uniqueNodes) == 1) {
146         res <- uniqueNodes[[1]]
147         return(res)
148     }
149     else { ## else length(uniqueNodes > 1)
150         ancestors <- lapply(nodes, OMPancestors, phy=phy, type="ALL")
151         res <- getNode(phy, max(Reduce(intersect, ancestors)))
152         return(res)
153     }
154 }#OMPMRCA
155
156 #####
157 # shortestPath
158 #####
159 OMPshortestPath <- function(phy, node1, node2){
160     ## conversion from phylo, phylo4 and phylo4d
161     if (class(phy) == "phylo4d") {
162         x <- extractTree(phy)
163     }
164     else if (class(phy) != "phylo4"){
165         x <- as(phy, "phylo4")
166     }
167
168     ## some checks
169     ## if (is.character(checkval <- checkPhylo4(x))) stop(checkval) # no need
170     t1 <- getNode(x, node1)
171     t2 <- getNode(x, node2)
172     if(any(is.na(c(t1,t2)))) stop("wrong node specified")

```

```

173     if(t1==t2) return(NULL)
174
175     ## main computations
176     comAnc <- OMPMRCA(x, t1, t2) # common ancestor
177     desComAnc <- OMPdescendants(x, comAnc, type="all")
178     ancT1 <- OMPancestors(x, t1, type="all")
179     path1 <- intersect(desComAnc, ancT1) # path: common anc -> t1
180
181     ancT2 <- OMPancestors(x, t2, type="all")
182     path2 <- intersect(desComAnc, ancT2) # path: common anc -> t2
183
184     res <- union(path1, path2) # union of the path
185     ## add the common ancestor if it differs from t1 or t2
186     if(!comAnc %in% c(t1, t2)){
187         res <- c(comAnc, res)
188     }
189
190     res <- getNode(x, res)
191
192     return(res)
193 } # end shortestPathe(isDes) <- "logical"
194
195     ## for internal nodes only, drop self (not sure why this rule?)
196     int.node <- intersect(node, nodeId(phy, "internal"))
197     isDes[cbind(match(int.node, descendant),
198               match(int.node, node))] <- FALSE
199
200     ## if only tips desired, drop internal nodes
201     if (type=="tips") {
202         isDes[descendant %in% nodeId(phy, "internal"),] <- FALSE
203     }
204
205     res <- getNode(phy, descendant[isDes[, seq_along(node)]]))
206 }
207
208     res
209 }#OMPdescendants
210
211 ## get ancestors (all nodes)
212 OMPancestors <- function(phy, node, type=c("all", "parent", "ALL")) {
213     type <- match.arg(type)
214
215     ## look up nodes, warning about and excluding invalid nodes
216     oNode <- node
217     node <- getNode(phy, node, missing="warn")
218     isValid <- !is.na(node)
219     node <- as.integer(node[isValid])
220
221     if (length(node) == 0) {
222         return(NA)
223     }
224
225     if (type == "parent") {
226         res <- lapply(node, function(x) ancestor(phy, x))
227     } else {
228         ## edge matrix must be in postorder for the C function!
229         if (phy@order=="postorder") {
230             edge <- phy@edge
231             } else {

```

```

232     edge <- reorder(phy, order="postorder")@edge
233   }
234   ## extract edge columns
235   ancestor <- as.integer(edge[, 1])
236   descendant <- as.integer(edge[, 2])
237
238   ## TODO: REPLACE C call with OMP implementation of ancestors
239   ## return indicator matrix of ALL ancestors (including self)
240   isAnc <- .Call("OMPancestors", node, ancestor, descendant)
241   storage.mode(isAnc) <- "logical"
242
243   ## drop self if needed
244   if (type=="all") {
245     isAnc[cbind(match(node, descendant), seq_along(node))] <- FALSE
246   }
247   res <- lapply(seq_along(node), function(n) getNode(phy,
248     descendant[isAnc[,n]]))
249 }
250 names(res) <- as.character(oNode[isValid])
251
252 ## if just a single node, return as a single vector
253 if (length(res)==1) res <- res[[1]]
254 res
255 }#OMPancestors
256
257 OMPMRCA <- function(phy, ...) {
258   nodes <- list(...)
259   ## if length==1 and first element is a vector,
260   ## use it as the list
261   if (length(nodes)==1 && length(nodes[[1]])>1) {
262     nodes <- as.list(nodes[[1]])
263   }
264
265   ## Correct behavior when the root is part of the nodes
266   testNodes <- lapply(nodes, getNode, x=phy)
267   ## BMB: why lapply, not sapply?
268   lNodes <- unlist(testNodes)
269   if (any(is.na(lNodes)))
270     stop("nodes not found in tree: ", paste(names(lNodes)[is.na(lNodes)],
271       collapse=", "))
272   uniqueNodes <- unique(testNodes)
273   root <- nTips(phy)+1
274   ## Handles case where root is a node of interest, return root
275   if(root %in% uniqueNodes) {
276     res <- getNode(phy, root)
277     return(res)
278   }
279   ## Correct behavior in case of MRCA of identical taxa
280   if(length(uniqueNodes) == 1) {
281     res <- uniqueNodes[[1]]
282     return(res)
283   }
284   else { ## else length(uniqueNodes) > 1)
285     ancests <- lapply(nodes, OMPancestors, phy=phy, type="ALL")
286     res <- getNode(phy, max(Reduce(intersect, ancests)))
287     return(res)
288   }
289 }#OMPMRCA
290

```

```

291
292 #####
293 # shortestPath
294 #####
295 OMPshortestPath <- function(phy, node1, node2){
296   ## conversion from phylo, phylo4 and phylo4d
297   if (class(phy) == "phylo4d") {
298     x <- extractTree(phy)
299   }
300   else if (class(phy) != "phylo4"){
301     x <- as(phy, "phylo4")
302   }
303
304   ## some checks
305   ## if (is.character(checkval <- checkPhylo4(x))) stop(checkval) # no need
306   t1 <- getNode(x, node1)
307   t2 <- getNode(x, node2)
308   if(any(is.na(c(t1,t2)))) stop("wrong node specified")
309   if(t1==t2) return(NULL)
310
311   ## main computations
312   comAnc <- OMPMRCA(x, t1, t2) # common ancestor
313   desComAnc <- OMPdescendants(x, comAnc, type="all")
314   ancT1 <- OMPancestors(x, t1, type="all")
315   path1 <- intersect(desComAnc, ancT1) # path: common anc -> t1
316
317   ancT2 <- OMPancestors(x, t2, type="all")
318   path2 <- intersect(desComAnc, ancT2) # path: common anc -> t2
319
320   res <- union(path1, path2) # union of the path
321   ## add the common ancestor if it differs from t1 or t2
322   if(!comAnc %in% c(t1,t2)){
323     res <- c(comAnc, res)
324   }
325
326   res <- getNode(x, res)
327
328   return(res)
329 } # end shortestPath

```

A.3 CUDA Code

Here are some highlights of the coming CUDA code.

Line 7: Node struct contains information about tree's nodes (node id, ancestor id, node label)

Line 13: setNode function translates R tree nodes to node structs

Line 22: kernel function GPU function that computes ancestors of all nodes in the tree

Line 67: shortestPath function takes an array of nodes, number of nodes given, and two character array labels and finds the shortest path between the given nodes. An error message is printed if the given nodes are invalid input (ex: they are a parent-child)

Line 196: main function driver function, reads in input and calls the shortestPath function

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <cuda.h>
4
5 //to compile: nvcc shortestPath.cu
6
7 struct node
8 {
9     int nodeID, ancestor;
10    char label[20]; //max size of label is 20
11 }; //node
12
13 void setNode(node &phy, int numNodes, int id, int aID, char * label)
14 {
15     phy.nodeID = id;
16     phy.ancestor = aID;
17     memset(phy.label, '\0', sizeof(label));
18     strcpy(phy.label, label);
19 } //setNode
20
21
22 --global-- void kernel(node * array, int numNodes, int id1, int id2,
23                        int * ancestorID1, int * ancestorID2)
24 {
25     int idx = blockIdx.x * blockDim.x + threadIdx.x;
26     if (idx < numNodes)
27     {
28         if (array[idx].nodeID == id1) //if found target node
29         {
30             int ancestorIndex = 0;
31             node temp = array[idx]; //start from current node
32             while (temp.ancestor != 0)
33             {
34                 ancestorID1[ancestorIndex++] = temp.ancestor; //add all ancestors
35                 for (int i=0; i<numNodes; i++)
36                 {
37                     if (array[i].nodeID == temp.ancestor)
38                     {
39                         temp = array[i];
40                         break;
41                     } //if
42                 } //for
43             } //while
44         } //if
45         else if (array[idx].nodeID == id2) //if found target node
46         {
47             int ancestorIndex = 0;
48             node temp = array[idx]; //start from current node
49             while (temp.ancestor != 0)
50             {
51                 ancestorID2[ancestorIndex++] = temp.ancestor; //add all ancestors
52                 for (int i=0; i<numNodes; i++)
53                 {
54                     if (array[i].nodeID == temp.ancestor)
55                     {
56                         temp = array[i];
57                         break;
58                     } //if
59                 } //for
60             } //while
61         } //if
62     } //if
63 }
```



```

64 } //kernel
65
66
67 void shortestPath(node * phy, int numNodes, char * label1, char * label2)
68 {
69     node * deviceArray;
70     int * deviceID1;
71     int * deviceID2;
72     int * ancestorID1 = new int[numNodes]; //initialize max size to number of nodes
73     int * ancestorID2 = new int[numNodes];
74     float blockSize = 1024; //num threads per block
75
76     //check if invalid query
77     node temp1, temp2;
78     for (int i=0; i<numNodes; i++)
79     {
80         ancestorID1[i] = 0;
81         ancestorID2[i] = 0;
82         if (strcmp(label1, phy[i].label) == 0)
83             temp1 = phy[i];
84         else if (strcmp(label2, phy[i].label) == 0)
85             temp2 = phy[i];
86     } //for
87
88     if ((temp1.ancestor == temp2.nodeID) || (temp2.ancestor == temp1.nodeID))
89     {
90         printf("named integer(0)\n");
91         return;
92     } //if
93
94     //allocate device memory
95     cudaMalloc(&deviceArray, sizeof(node) * numNodes);
96     cudaMalloc(&deviceID1, sizeof(int) * numNodes);
97     cudaMalloc(&deviceID2, sizeof(int) * numNodes);
98     cudaMemcpy(deviceArray, phy, sizeof(node) * numNodes, cudaMemcpyHostToDevice);
99     cudaMemcpy(deviceID1, ancestorID1, sizeof(int) * numNodes, cudaMemcpyHostToDevice);
100    cudaMemcpy(deviceID2, ancestorID2, sizeof(int) * numNodes, cudaMemcpyHostToDevice);
101
102    dim3 dimBlock(blockSize);
103    dim3 dimGrid(ceil(numNodes/blockSize));
104
105    //compute ancestors
106    kernel <<< dimGrid, dimBlock >>> (deviceArray, numNodes, temp1.nodeID, temp2.nodeID, ↵
        deviceID1, deviceID2);
107    cudaMemcpy(ancestorID1, deviceID1, sizeof(int) * numNodes, cudaMemcpyDeviceToHost);
108    cudaMemcpy(ancestorID2, deviceID2, sizeof(int) * numNodes, cudaMemcpyDeviceToHost);
109    cudaFree(deviceArray);
110    cudaFree(deviceID1);
111    cudaFree(deviceID2);
112
113    //find shortest path
114    int * path = new int[numNodes];
115    int currentPath = ancestorID1[0];
116    int pathIndex = 0;
117    bool isLCAPath = false;
118    //check if path converges at LCA
119    for (int i=0; i<numNodes; i++)
120    {
121        path[i] = 0;
122        if (temp1.nodeID == ancestorID2[i])
123        {
124            for (int j=0; j<i; j++)
125                path[j] = ancestorID2[j];
126            isLCAPath = true;
127            break;
128        } //if
129        else if (temp2.nodeID == ancestorID1[i])
130        {

```

```

131         for (int j=0; j<i; j++)
132             path[j] = ancestorID1[j];
133         isLCAPath = true;
134         break;
135     } //else if
136 } //for
137
138 //if one node is the ancestor of another
139 if (!isLCAPath)
140 {
141     for (int i=0; i<numNodes; i++)
142     {
143         for (int j=0; j<numNodes; j++)
144         {
145             if (currentPath == ancestorID2[j])
146                 break;
147             if ((ancestorID2[j] == 0) || (j == numNodes-1))
148             {
149                 path[pathIndex++] = ancestorID1[i];
150                 currentPath = ancestorID1[i];
151                 break;
152             } //if
153         } //for
154     } //for
155
156     if (pathIndex == 0)
157         path[pathIndex++] = currentPath;
158
159     for (int i=0; i<numNodes; i++)
160     {
161         if (ancestorID2[i] == currentPath)
162             break;
163         path[pathIndex++] = ancestorID2[i];
164     } //for
165 } //if
166
167 for (int i=0; i<numNodes; i++)
168 {
169     if (path[i] == 0)
170         break;
171     for (int j=0; j<numNodes; j++)
172     {
173         if (path[i] == phy[j].nodeID)
174         {
175             printf("%s ", phy[j].label);
176             break;
177         } //if
178     } //for
179 } //for
180
181 printf("\n");
182 for (int i=0; i<numNodes; i++)
183 {
184     if (path[i] == 0)
185         break;
186     printf("%d ", path[i]);
187 } //for
188 printf("\n");
189
190 delete [] ancestorID1;
191 delete [] ancestorID2;
192 delete [] path;
193 } //shortestPath
194
195
196 int main()
197 {
198     int numNodes = 27;

```

```

199     node * phy = new node[numNodes];
200     FILE * infile = fopen("geospiza", "r");
201
202     int nodeID, ancestor;
203     char label[20];
204     for (int i=0; i<numNodes; i++)
205     {
206         fscanf(infile, "%d", &nodeID);
207         fscanf(infile, "%d", &ancestor);
208         fscanf(infile, "%s", &label);
209         setNode(phy[i], numNodes, nodeID, ancestor, label);
210     } //for
211     fclose(infile);
212
213 //test shortest path
214     shortestPath(phy, numNodes, "fusca", "fortis");
215
216     delete [] phy;
217     return 0;
218 } //main

```

B Who Did What

Alicia wrote the OpenMP implementation. Bryan wrote the CUDA implemenation. Raymond wrote the RSnw implementation. We worked on running tests and writing the report in L^AT_EX.

References

- [1] F. Michonneau, <http://cran.r-project.org/web/packages/phylobase/phylobase.pdf> February 20, 2015.